

# MSDS 7333 Fall 2020: Case Study 2

Jayson Barker, Brandon Croom, Shane Winestock

## Business Understanding

The Cherry Blossom Ten Mile Run is an annual 10-mile race in Washington, D.C. The race was originally founded in 1973 and has been run continuously ever since. The race is meant to coincide with the bloom of the cherry blossoms in the DC area and is part of the National Cherry Blossom Festival. The race started as a training race for professional runners planning to compete in the Boston Marathon. Over the years the race has evolved into a race that attracts both professional and local runners of all levels. Evaluating this change in runners is a critical component for the race planners to understand how to position and market the race to attract participants.

This report focuses on the 1999 through 2012 years of the race. The goal is to specifically analyze female runners over the years to determine the answers to a few key questions that the race planners need for marketing purposes. Specifically the following questions will be addressed:

- Have age distributions changed over time?
- Have race times increased or decreased over time?
- What trends should the planners be aware of that may occur in future years?

## Data Acquisition

In order to perform the analysis the data for races from 1999 through 2012 must be acquired. The easiest way to gather this data is through web scraping of the Cherry Blossom race website. The website contains web pages listing all of the race results through the specified time period.

Web scraping has many aspects that must be taken into account when performed. First and foremost, the data must be allowed to be scraped from websites. The content of a website may be a revenue generation opportunity or built out content that the website owner has spent considerable time curating. When acquiring data from websites it must first be verified that the website allows scraping of the site's data. Scraping websites without site owner's consent may result in litigation if the data is presented without attribution and may be in violation of copyright laws. In this particular case, the data from the Cherry Blossom site is being presented to the race planners. There are no legal concerns about scraping the Cherry Blossom site data in this context since the data is being presented back to the race planners.

A second aspect that must be taken into account when scraping websites is understanding the data availability and working through the differences that may occur overtime with individual site webpages. In the Cherry Blossom case, the race results were housed on individual webpages for each race year. However, each year's webpages has subtle differences that required changes in the acquisition method to ensure the data was captured correctly. For example, some years of the Cherry Blossom race captured only race completion time, while others included completion time and gun time. The web scraping must be adapted to take all of these possible changes into account.

In building out the specific web scraping capability leveraged for this analysis a few key items in the individual webpages needed to be addressed. Each webpage for each year is analyzed for key items that must be taken

into account to ensure the data can be scraped in a consistent fashion. Each webpage uses equal marks ('=') as column header delimiters. The specific spacing for each column for each year must be accounted for to ensure the correct fields were consistently mapped. In addition the necessary webpage tags (the layout information for each individual webpage) much be searched to find the correct tag that contains the race results. The years 1999, 2000 and 2009 had different formatting for these tags than other years, thus special analysis was required to scrape these years to ensure consistent results.

The results of scraping the Cherry Blossom site for the women's results provides us the following record counts for the years in question. This information provides an initial check that the web scraping is working as expected and provides an individual record set by year for initial data cleansing.

```
##      Year Count
##  1: 1999  2356
##  2: 2000  2167
##  3: 2001  2973
##  4: 2002  3335
##  5: 2003  3544
##  6: 2004  3903
##  7: 2005  4338
##  8: 2006  5441
##  9: 2007  5696
## 10: 2008  6401
## 11: 2009  8329
## 12: 2010  8859
## 13: 2011  9034
## 14: 2012  9733
```

## Data Cleaning

Now that the data has been broken down into individual record sets for each year, an initial look at the data structure for each individual record set is warranted. This will ensure that data mapped into the individual data sets represents the correct data types for the values that are expected to be pulled from the race webpages. Since all individual record sets are pulled in the same format evaluating the 2000 record set will provide a quick indication that data was mapped correctly.

```
##      Place      DivTot      Num      Name      Ag      Hometown
## "character" "character" "character" "character" "character" "character"
##           Gun          Net
## "character" "character"
```

From the above information it can be seen that every field is being read as a character field. Given that the data is being scraped from the web, this makes sense. Data conversion will be necessary to the respective data types as individual record sets are combined into a full record set.

The evaluation of missing data is another items that is best addressed at the individual record set level. Performing missing data checks at this level makes addressing them much easier prior to combining the data into a full record data set. The table below shows the results of missing data analysis for each of the individual record sets.

```
##      Race Year Missing Record Count
##  1:      1999           0
##  2:      2000           0
```

```
## 3:      2001      0
## 4:      2002      0
## 5:      2003      0
## 6:      2004      0
## 7:      2005      0
## 8:      2006      0
## 9:      2007      0
## 10:     2008      0
## 11:     2009      0
## 12:     2010      0
## 13:     2011      0
## 14:     2012      0
```

From the table above, there is no data missing in the records. This shows that the data entered in to the web site is complete and that the data scraping methodology worked as expected.

As previously noted, the individual record sets have all data defined as characters. The time values (Gun, Net) and the age values need to be converted to numeric data types to ensure further analysis is correct. At the same time all of the data needs to be merged into a single record set containing all years. As part of this merging, a new data field for race time (Net\_Conv) will be created. This data field will contain the race time converted to minutes. Just as with the individual record sets the evaluation of data types and missing values for the combined record set will occur to ensure data is ready for analysis.

Checking the combined record set data types, as noted below, all values are now more indicative of the respective data column. Age (Ag) is now numeric and the race time (Net\_Conv) is also numeric.

```
##      Year      Place      DivTot      Name      Ag      Hometown
## "character" "character" "character" "character"  "numeric" "character"
##      Net      Pace      Net_Conv
## "character" "character"  "numeric"
```

Evaluating the data for missing values produces the following results:

```
##      Missing Record Count
## 1:      0
```

Based on the missing value check, it can be seen that the data has remained intact throughout the transformations. One additional check to provide summary statistics of the data will be performed prior to beginning analysis. This summary statistics check provides another sanity check for data evaluation.

```
##      Year      Place      DivTot      Name
## Length:75839  Length:75839  Length:75839  Length:75839
## Class :character  Class :character  Class :character  Class :character
## Mode  :character  Mode  :character  Mode  :character  Mode  :character
##
##
##      Ag      Hometown      Net      Pace
## Min.   : 0.00  Length:75839  Length:75839  Length:75839
## 1st Qu.:26.00  Class :character  Class :character  Class :character
## Median :31.00  Mode  :character  Mode  :character  Mode  :character
## Mean   :32.73
## 3rd Qu.:39.00
```

```
## Max.      :87.00
##      Net_Conv
## Min.      : 51.73
## 1st Qu.   : 89.00
## Median    : 98.18
## Mean      : 99.31
## 3rd Qu.   :108.40
## Max.      :170.97
```

The summary statistics above reconfirm the information that was expected. The character fields should not have any statistical values. Additionally, the Age (Ag) and race time (Net\_Conv) fields do show reasonable statistical values. It should be noted that the Age (Ag) field does show a minimum value of 0.00. This indicates that there may be incorrect data input into the age field for some runners or that the data was omitted. Evaluating the combined record set to pick out zero age record results in the data shown below.

```
##      Year Place DivTot      Name Ag      Hometown      Net Pace
## 7129 2001  2611      0 Loretta CUCE      0 Alexandria VA      1:53:38      0
##      Net_Conv
## 7129 113.6333
```


Evaluation of the race results web page for the 2001 race does confirm that the age entered for this individual was 0. As such this should be treated as an invalid data point, with respect to age. On a positive note, there is only one record in this data set with a zero value for age.

## Data Analysis

Now that the data has been scraped and cleaned for analysis, the process of answering the requested business questions can begin. The initial analysis will start with an analysis of racers ages over time. It will then move into the evaluation of race completion times over time.

## Racer Age Analysis


Women's Racers Ages – All Years



From the box plot it can be seen that 2006 had a lot of younger racers compared to the other years. Additionally, in 2006 many of these younger racers were at the tail end of the race times. This could indicate an influx of non-professional runners during the race for that year.

Diving into the distribution of ages across the race years the plot below shows a bimodal distribution of ages. This distribution reaffirms the large number of younger racers seen in 2006.

## Women's Racers Age Distribution – All Years



Continuing to evaluate runners ages the data provides an opportunity to evaluate mean age by race year as shown in the table below.

##	Year	Mean Age
## 1:	1999	34.9009353741497
## 2:	2000	35.5535549399815
## 3:	2001	34.8061911170929
## 4:	2002	35.1378378378378
## 5:	2003	35.0513833992095
## 6:	2004	33.9296534017972
## 7:	2005	34.1692485549133
## 8:	2006	18.3847994111152
## 9:	2007	33.4678571428571
## 10:	2008	33.2102548069408
## 11:	2009	33.0775147217882
## 12:	2010	33.0376171964327
## 13:	2011	33.7369878183832
## 14:	2012	33.8779936272998
## 15:	Total	32.7250754888646


This table confirms the distribution plot and box plots above with respect to racers in 2006 skewing much younger for that year. Using the mean age of all races as 32, looking specifically at 2006 it can be seen that of the 5435 racers in 2006, 3784 of them were below the mean race age.

Digging deeper into the 2006 data specifically, the distribution plot below shows that a large portion of the runners in 2006 were under 18 years of age. There could be many factors that account for this lower age skew.


- Runners could be truly entering at a younger age
- Runners could be providing lower ages
- Data could have been entered incorrectly

Each of these cases could take additional data and understanding of the circumstances to come to a concrete conclusion.

### Women's Racers Age Distribution – 2006




Reevaluating the age data to account for the 2006 outlier (filtering for racers that are older than 18 years of age) the following distribution can be seen:



Accounting for the extremely young outliers in the distribution above it can be seen that there is still some skew towards younger ages. The outliers in this case are 2002 and 2007 where ages were actually a bit older. Leveraging the same data in the scatter plot below it can be seen that there is a stronger concentration of ages in the 2009-2012 year range compared to previous years. This could indicate that more racers are entering the race. This can be reaffirmed by looking again at the number of racers in each yearly race.





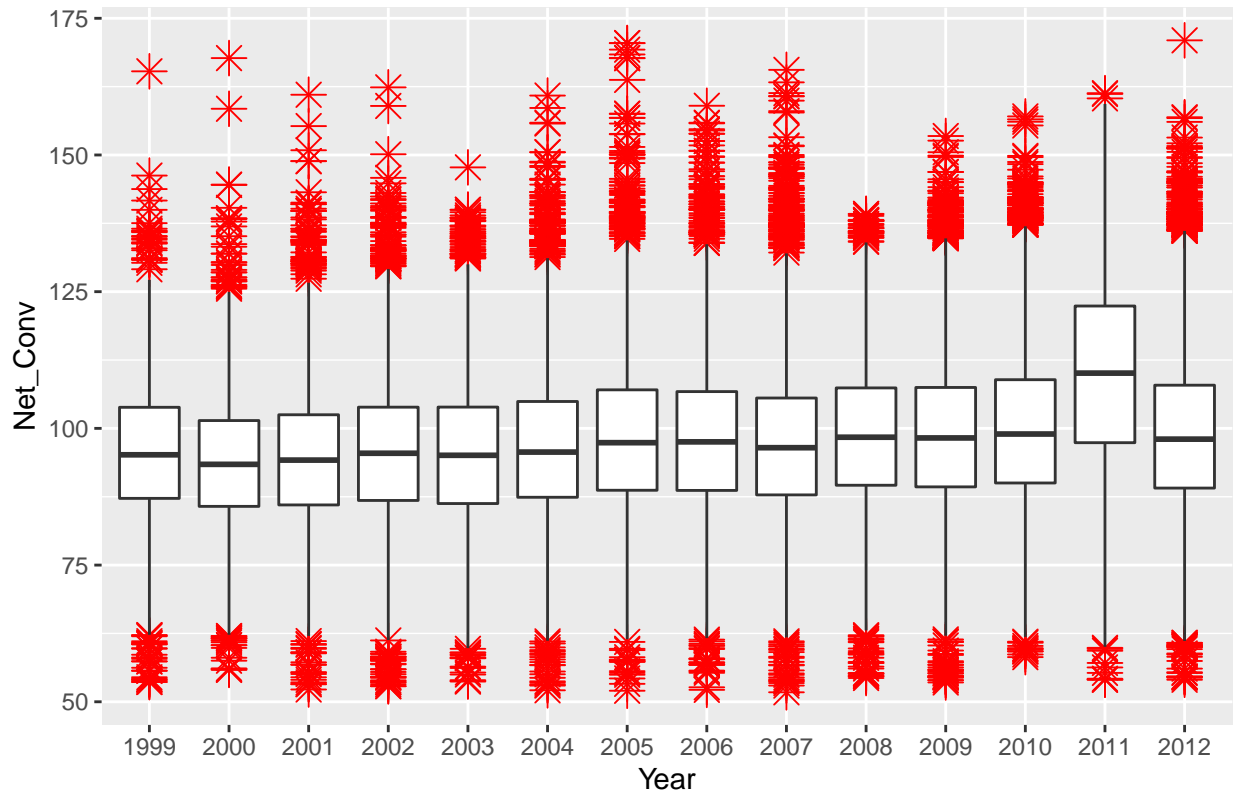
##	Year	Count
## 1:	1999	2356
## 2:	2000	2167
## 3:	2001	2973
## 4:	2002	3335
## 5:	2003	3544
## 6:	2004	3903
## 7:	2005	4338
## 8:	2006	5441
## 9:	2007	5696
## 10:	2008	6401
## 11:	2009	8329
## 12:	2010	8859
## 13:	2011	9034
## 14:	2012	9733

Both the scatterplot and table above confirm that the number of racers has continued to increase from 1999 through 2012.

## Race Time Analysis

Just as with the age data there are interesting trends that emerge when evaluating race time data. A simple box plot of race times begins to show some of these points of interest. In the box plot below of Net Race times over all years it can be seen that on average the race times were consistent in all years except for 2011.

Women's Racers Net Times – All Years



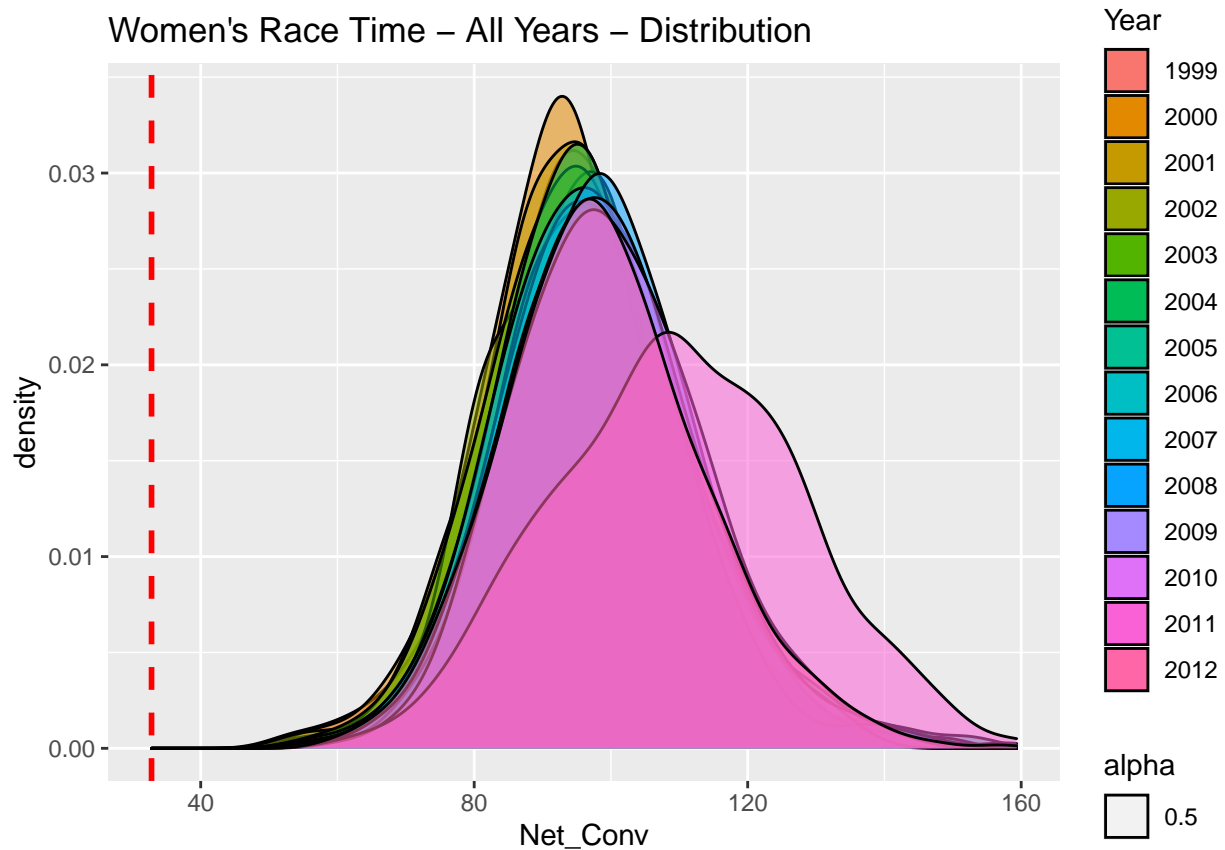
Digging deeper into mean runner time provides the following table:

##	Year	Mean Time
## 1:	1999	95.538102324263
## 2:	2000	93.9138504155125
## 3:	2001	94.6531123822342
## 4:	2002	95.9373323323323
## 5:	2003	95.5493459439112
## 6:	2004	96.7006718014549
## 7:	2005	98.6066319845857
## 8:	2006	98.1578241933505
## 9:	2007	97.1735476190476
## 10:	2008	98.7727710906154
## 11:	2009	98.8117333653808
## 12:	2010	99.92128401555
## 13:	2011	109.904671465485
## 14:	2012	99.0223027375201
## 15:	Total	99.313335047491

From the mean time table above the 2011 outlier is reaffirmed. This outlier is almost 10 minutes higher than the mean of all races. Additional data would be needed to evaluate possible causes of this outlier discrepancy. Multiple factors could be at play to cause this outlier. Some examples of possible impacts could be: racers could have truly run the race at a slower pace, weather conditions could have impacted racers, or even the total number of racers could have impacted these times. Based on the data available, 2011 had the second highest total number of racers, with 9034, however this was only an increase of approximately 200 racers from 2010. In addition, the number of racers in 2012 went up to 9733 which was a much larger increase

and race times for 2012 were back around the mean. Additional analysis would need to be performed with additional data sets to truly understand what caused 2011 to be such an outlier in race time.

The distribution chart below also reaffirms that 2011 was an outlier year.



## Conclusion

In conclusion the data on the Cherry Blossom race from 1999 through 2012 has provided some interesting insights. These insights should help race planners in evaluating any changes that may need to occur to maintain excitement about the race.

With respect to racers age, it was shown that the mean racers age of approximately 32 years old has remained consistent over time. There was a data anomaly in 2006 where many young racers performed in the race, however that anomaly was not seen in other years.

With respect to race completion times, like age, these too have remained fairly consistent over the 1999 through 2012 time period. The year 2011 was an outlier in this data and those results were not seen again in the data set. On average racers were finishing the the race in about 99 minutes.

The data also confirm that the number of racers entering the race continues to increase year over year. 1999 had 2365 runners while 2012 had 9733 runners. This is an extremely large increase. From a business perspective it could indicate that interest in the race still exists.

All in all, both the age and time data seem pretty consistent over the 13 year span of data that was evaluated. As the number of racers continue to increase over the years these values may shift as more racers begin to enter the race.

From a race planner perspective this data can assist in planning for future races on many levels. It can assist with planning out the race course. On average the racers are completing in about 99 minutes, in a congested city such as Washington, D.C., understanding how long streets may be blocked off and impact traffic around the city. In addition, this information may go into overall route planning by providing alternative race routes that fall into the length of the race but also can be completed in the average time. Understanding the average ages of the runners could largely help the planners in a marketing sense for the race. Knowing that the racer ages skew toward the low 30s provides the race planners with information about how best to reach those racers. It may be more beneficial for the race planners to market on websites and social media sites given the average racer age than to advertise in running magazines or newspapers.

Overall, the trends look to be relatively steady with respect to runners ages and completion times. There is a consistent increase in number of runners attending the races, however ages and finish times remain consistent. This trend looks to be consistent over the 13 period analyzed with only a few outliers in each case.